

PRESIDENCY UNIVERSITY

Presidency School of Artificial Intelligence and Advanced Computing

AI Centre of Excellence (Accelerated by NVIDIA)

Ref No: PU/PSAIAC/SIP/2026/20241CSN0042

Date: May 31, 2026

To,
Mr. / Ms. Adarsha p.g
Roll No: 20241CSN0042
Program / Specialization: CSN - Computer Science & Engineering (Networks)

Semester:5th Semester, Undergraduate
Presidency University, Bengaluru

Subject: Offer of Admission and Guidelines for the Intensive AI GPU Summer Internship Program (Summer 2026)

Dear Adarsha p.g,

Following the rigorous evaluation of your academic standing, motivation statement, we are pleased to offer you formal admission to the intensive, 2-week AI GPU Summer Internship Program (Summer 2026) at Presidency University. This elite training track is organized by the Presidency School of Artificial Intelligence and Advanced Computing in direct collaboration with our NVIDIA Accelerated AI Centre of Excellence. By accepting this offer, you will join a select cohort of students training on world-class computational architecture.

1. Individualized Training Schedule, Venue & Fee Allocations

Your program assignment contains specific technical parameters mapped to your ledger profile. You are required to attend sessions strictly within your allocated block:

Parameter	Individualized Allocation Detail
Assigned Cohort Phase	Slot 1 will be from June 1 to 12 (10 days)
Daily Time Slot Allocation	Slot A: 09:00 AM - 12:00 PM
Assigned Laboratory Venue	Lab (DGL01-DGL02) - D-Block GPU Computing Lab Cluster
Registration Fee Category	Paid
Program Structural Volume	3 Hours/day over 10 Instructional Days (Total: 30 Contact Hours)
Academic Credit Load	2 Transferable Academic Credits

2. Structural Curriculum & Timeline Overview

The program is structured as a 30-hour intensive workflow incorporating 20% Lecture/Theory (6 hrs), 20% Guided Coding Workshops (6 hrs), 33% GPU Hands-on Labs (10 hrs), 20% Mentored Project Work (6 hrs), and 7% Final Assessment and Presentations (2 hrs). The 10 instructional days advance through the following core technical modules:

- **Week 1 (AI Foundations & Deep Learning Essentials):** Orientation & Environment Setup (Day 1); Machine Learning Fundamentals & Accelerated Sklearn on GPU (Day 2); Neural Networks from Scratch & PyTorch Basics on H200 (Day 3); Convolutional Neural Networks (CNNs) & GPU Profiling with Nsight (Day 4); and Hands-on Google Antigravity & Foundation Quiz (Day 5).
- **Week 2 (Advanced AI, LLMs & Capstone Project):** Transformer Architecture Deep Dive (Day 6); LLM Fine-Tuning with LoRA & PEFT via Hugging Face on H200 (Day 7); Generative AI via Stable Diffusion XL pipelines (Day 8); GPU Optimization Optimization (Mixed Precision, FlashAttention, Tensor Parallelism) & Capstone Execution (Day 9); and Capstone Project Presentations, Peer Review & Closing Ceremony (Day 10).

3. Hardware Infrastructure Utilization & AI Ethics

Throughout this internship, you will have exclusive access to a dedicated container instance on our state-of-the-art **NVIDIA H200 Tensor Core GPU Cluster** via JupyterHub. The H200 architecture is built on the Hopper (GH100) architecture, featuring 141 GB of HBM3e memory, 4.8 TB/s memory bandwidth, 3,958 FP8 Tensor TFLOPS, and integrated Transformer Engines optimized for large language model workloads.

- **Compute Provision & Suspend Policy:** GPU instances run strictly during your scheduled lab hours and will auto-suspend post-session. Temporary local container storage is wiped daily; all code must be pushed to your persistent repository before session termination.
- **Prohibited Operations:** Utilizing the H200 infrastructure for cryptocurrency mining, network scanning, unauthorized model training, or sharing infrastructure credentials will result in immediate program termination and academic suspension.
- **AI Ethics & Integrity:** While the use of generative AI tools for code layout exploration is permitted, all AI-generated code snippets must be explicitly declared in your lab logs. Plagiarism or undocumented generation will cause automatic program disqualification.

4. Mandatory Rules, Regulations, and Attendance Framework

To maintain strict institutional compliance and structural credit equivalency, you must adhere to the following framework:

- **Mandatory Attendance:** Attendance across all 10 instructional days is strictly compulsory. Any student missing more than one (1) day of instruction without prior written administrative approval from the Program Chair will be automatically disqualified and will not receive a certificate or credit.
- **Punctuality Constraint:** Late arrivals exceeding thirty (30) minutes from your designated start time (**Slot A: 09:00 AM – 12:00 PM**) will be marked as an absolute absence for that calendar day. Repeated late arrivals may result in immediate program removal.
- **Format Definiteness:** This is an elite, fully in-person computational lab program; no remote or hybrid participation is permitted.
- **Professional Conduct:** Professional behavior, discipline, and absolute adherence to instructions provided by mentors and coordinators are mandatory throughout the internship period.

5. Continuous Academic Assessments and Capstone Deliverables

Your performance and final credit allocation will be determined based on continuous assessment metrics weighing a total of 100 points. Successful completion requires achieving passing thresholds across all components:

1. **Daily Lab Submissions (35% Weight):** Submission of 7 core GPU lab programming modules (including CUDA Hello World, ResNet-18 Image Classification, BERT Sentiment Analysis, LLAMA 3 8B 4-bit QLoRA Fine-Tuning, SDXL Image Generation, and Nsight Profiling Sprint) at the end of each designated lab day.
2. **Week 1 Class Test / Quiz (10% Weight):** A conceptual evaluation on Core AI/ML and mathematical foundations administered on Day 5.
3. **Capstone Project (50% Cumulative Weight):** Independent execution of a chosen engineering track (Computer Vision, NLP, Generative AI, or GPU Systems Optimization) on the H200 GPU node. Final project deliverables include:
 - Fully documented, reproducible GitHub Code Repository containing functional code (20%).
 - Live Runtime Demonstration & Presentation Panel on the H200 GPU on Day 10 Afternoon (15%).
 - A 4–6 page Formal Technical Architecture Report covering problem statement, methodology, implementation, and results submitted on Day 10 Morning (15%).
4. **Peer Review & Participation (5% Weight):** Evaluated via active collaboration and peer code review cycles on Day 10.

To confirm your seat in your allotted phase, you are required to complete the final processing via your assigned payment ledger link, after which your JupyterHub credentials will be securely dispatched to your registered university email. We congratulate you on your selection and look forward to watching you build the next generation of AI systems.

AI CoE Coordination Team:

Mr. Gyanesh Verma
 Mr. Vinayak Raju Kage
 Mr. Thatimakula Sai Kumar
 Ms. Varalakshmi T

Dr. Robin Rohit Vincent

Head, AI CoE (NVIDIA)

Presidency School of AI & Advanced Computing

STUDENT ACKNOWLEDGMENT & UNDERTAKING

I, **Adarsha p.g**, with **20241CSN0042** hereby acknowledge that I have carefully read, completely understood, and agree to strictly abide by the attendance benchmarks, strict computational infrastructure usage policies, academic integrity codes, assessment frameworks, and mandatory capstone deliverable structures (functional project, live runtime demonstration, 4-6 page technical report, and presentation slides) outlined in this offer letter. I understand that failure to comply with these stated institutional requirements will lead to immediate program disqualification, non-issuance of the Internship Completion Certificate, and non-award of the 2 associated academic credits.

Student Signature
Adarsha p.g

Date & Place

AI Centre of Excellence (Accelerated by NVIDIA)
Presidency School of Artificial Intelligence and Advanced Computing
PRESIDENCY UNIVERSITY